

OS: C Threading API (pthreads)

CIT 595
Spring 2010

Pthreads


- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)
- API specifies behavior of the thread library, implementation is up to development of the library
- Note: To compile program
 - gcc -lpthread filename.c
 - Must provide -lpthread flag as this library is dynamically linked (i.e. linked at runtime)

CIT 595

2

Designing Concurrent Programs with Pthreads

- Take advantage of Pthreads, the work should be organized into discrete, independent tasks which can execute concurrently
- For example, if routine1 and routine2 can be interchanged, interleaved and/or overlapped in real time, they are candidates for threading



CIT 595

Source: <https://computing.lnl.gov/tutorials/pthreads/>

3

Thread Creation

```
int pthread_create (pthread_t * thread, pthread_attr_t * attr, void *  
(*start_routine) (void *), void * arg)
```

- **thread**: unique identifier for the new thread returned by the subroutine
- **attr**: attribute object that may be used to set thread attribute
 - You can specify a thread attributes object, or NULL for the default values
 - Attributes such as stack size, priority, joinable or detachable
- **start_routine**: the C routine that the thread will execute once it is created
 - function that thread perform must be void * funcname (void *)
- **arg**: A single argument that may be passed to **start_routine**.
 - It must be passed by reference as a pointer cast of type void
 - NULL may be used if no argument is to be passed
 - It should be cast to its correct type in the function
- returns 0 when successful

CIT 595

4

Example thread1.c

```

struct threadData{
    int id;
};
typedef struct threadData thData;

int main(){
    pthread_t threads[NUM_THREADS];
    thData thdataArray[NUM_THREADS];
    int rc, t;
    for(t=0; t < NUM_THREADS; t++){
        thdataArray[t].id = t;
        printf("In main: creating thread %d\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello,
            (void *)&thdataArray[t]);
        if (rc){
            printf("ERROR: return code from pthread_create() is
                %d\n", rc);
            exit(-1);
        }
    }
}
//contd next slide

```

CIT 595

5

Example thread1.c contd ..

```

//Main contd
for (t = 0 ; t < NUM_THREADS ; t++) {
    rc = pthread_join( threads[t], NULL );
    if (rc != 0) {
        printf( "Error joining thread %d\n",
            thdataArray[t].id );
        exit(-1);
    }
}
return 0;
}

void *PrintHello(void *threadid){
    thData * temp = (thData *)threadid;
    printf("Hello World! It's me, thread
        #%d!\n", temp->id);
    pthread_exit(NULL);
}

```

■ Sample Output

```

In main: creating thread 0
In main: creating thread 1
In main: creating thread 2
In main: creating thread 3
In main: creating thread 4
In main: creating thread 5
In main: creating thread 6
Hello World! It's me, thread #0!
Hello World! It's me, thread #1!
Hello World! It's me, thread #6!
In main: creating thread 7
In main: creating thread 8
Hello World! It's me, thread #2!
Hello World! It's me, thread #5!
Hello World! It's me, thread #4!
Hello World! It's me, thread #7!
In main: creating thread 9
Hello World! It's me, thread #3!
Hello World! It's me, thread #8!
Hello World! It's me, thread #9!

```

CIT 595

6

Incorrect argument passing

- The code below passes the *address* of variable t, which is shared memory space and visible to all threads
 - As the loop iterates, the value of this memory location changes, possibly before the created threads can access it

```

int rc; int t;
for(t=0; t<NUM_THREADS; t++) {
    printf("Creating thread %ld\n", t);
    rc = pthread_create(&threads[t], NULL, PrintHello, (void *) &t);
    ...
}

```

CIT 595

7

Thread Termination

- void pthread_exit(void *value_ptr);
- *value_ptr* is a pointer to the object (variable, array, structure) returned by the thread
 - The value must not be of local scope otherwise it won't exist after the thread is destroyed
 - This value can be available to another thread in the same process
 - Can be NULL if not returning anything

CIT 595

8

Thread join

- One way to accomplish synchronization between threads
- Causes the calling thread to wait for another thread to terminate
 - For threads it important as we run the risk of executing an exit (reach end of main) which will terminate the process and all threads before the threads have completed
- `int pthread_join(pthread_t thread, void ** value_ptr)`
 - `thread` is the thread to wait on
 - `value_ptr` is the value given to `pthread_exit()` by the terminating thread
 - returns 0 to indicate success
- Usage

```
void * return_val;
....
//after code creating threads
if(pthread_join(worker_thread, &return_val)){
    printf("Error while waiting on thread\n");
    exit(1);
}
```

CIT 595

9

Mutex

- Mutex is short for mutual exclusion
- Primary means of implementing thread synchronization and for protecting shared data when multiple writes occur
 - Provides lock mechanism for shared data
- To create a mutex:
`int pthread_mutex_init(pthread_mutex_t *mutex, pthread_mutexattr_t * attr)`
 - `mutex` is the lock (of type `pthread_mutex_t`)
 - `attr` is the lock attributes
 - NULL by default
- To lock: `int pthread_mutex_lock(pthread_mutex_t *mutex)`
 - If lock is already locked the calling thread is blocked
 - If lock is not locked the calling thread acquires it
 - returns 0 on success
- To Unlock: `pthread_mutex_unlock`

CIT 595

10

Typical sequence when using a mutex

- Create and initialize a mutex variable
- Several threads attempt to lock the mutex
- Only one succeeds and that thread owns the mutex
- The owner thread performs some set of actions
- The owner unlocks the mutex
- Another thread acquires the mutex and repeats the process
- Finally the mutex is destroyed

CIT 595

11

Mutex (contd..)

- Declare variables as globals i.e. outside all methods
 - E.g. `pthread_mutex_t myMutex;`
- For synchronization:

```
pthread_mutex_lock(&myMutex)
//critical section code
pthread_mutex_unlock(&myMutex)
```
- After all work is done, need to destroy them
 - `pthread_mutex_destroy (&myMutex);`
- See example `thread2.c`

CIT 595

12

Deadlock in Resource Sharing Environment

- A **deadlock** occurs when 2 or more tasks (processes/threads) permanently block each other by each having a lock on a resource which the other tasks are trying to lock
- Deadlock can occur due to
 - Locks: Waiting to acquire locks on resources, such as objects, pages etc.
 - Sharing resources such as I/O devices printer, disks etc.

CIT 595

13

Example

- 2 Threads access 2 shared variables A and B
- Variable A is protected by lock x and variable B by lock y
- Here's what Thread 1 and Thread 2 need to do:

Thread 1

```
A = A + 10
B = B + 20
A = A + B
A = A + 30
```

Thread 2

```
B = B + 10
A = A + 20
A = A + B
B = B + 30
```

- Each must acquire locks for A and B
- Lets look at one way to do this

CIT 595

14

Approach

Thread 1

```
Lock(x)
A = A + 10
Lock(y)
B = B + 20
A = A + B
Unlock(y)
A = A + 30
Unlock(x)
```

Thread 2

```
Lock(y)
B = B + 10
Lock(x)
A = A + 20
A = A + B
Unlock(x)
B = B + 30
Unlock(y)
```

- Can we see a problem with this approach?
- How can we avoid the problem?

CIT 595

15

Cond and Wait

- Another means of synchronization
- Condition variables allow threads to synchronize based upon the actual value of data
 - If we want one thread to signal an event to another we need to use Conditional variables
 - pthread_cond_t *condVariable*
- Idea is one thread wait until a certain condition is true
 - Test condition
 - If not true, calls pthread_cond_wait(..) to block until it is
- Another thread makes the condition true and call pthread_cond_signal(...) to unblock the thread waiting
- To avoid race conditions, the conditional variable must use a mutex

CIT 595

16

Example

- Signalling thread

```
pthread_mutex_lock(&mutex);
flag = 1;
pthread_cond_signal(&condition);
pthread_mutex_unlock(&mutex);
```
- condition is conditional variable
 - type `pthread_cond_t`
- mutex is mutex variable
 - type `pthread_mutex_t`
- Waiting thread

```
pthread_mutex_lock(&mutex);
if(flag == 0)
    pthread_cond_wait(&condition, &mutex);
pthread_mutex_unlock(&mutex)
```
- Wait will automatically release the mutex while it waits
- After signal is received and thread is awakened, *mutex* will be automatically locked for use by the thread.
- Programmer is then responsible for unlocking *mutex* when the thread is finished with it

CIT 595

17

Producer Consumer Example

- Single producer thread, single consumer thread
 - Single shared buffer between producer and consumer
 - E.g. A fixed-size queue of print requests
 - One thread produces information – adds a print request to the queue
 - Other thread consumes information – takes a print request and prints it
- Condition and Wait
 - Involves mutual exclusion between producer and consumer
 - Due to use of same buffer
 - After producing an item, a producer should signal the consumer
 - After consuming, consumer should signal the producer
 - See example `thread3.c` for illustration

CIT 595

18

Misc.

- `pthread_self()` returns the unique, system assigned thread ID of the calling thread
- Thread cancelation
 - `int pthread_cancel(pthread_t thread)`
 - Causes the thread to be canceled (or terminated)
- Thread Attributes
 - By default, a thread is created with certain attributes
 - `pthread_attr_t` and `pthread_attr_t` are used to initialize/destroy the thread attribute object
 - `pthread_attr_getXXX(..)` gets the attribute and `pthread_attr_setXXX(..)` sets the attribute
- Mutex Attributes
 - Like threads, mutexes also have attributes
 - Related to thread scheduling (more details in scheduling topics)

CIT 595

19